InterviewForge AI — Public Technical Notes

InterviewForge AI uses ChromaDB as its vector database.
Its local large language model is Llama 3.2, served through Ollama.
The backend is written in Python with FastAPI; the frontend uses React.
The Knowledge Vault ingests PDF files, chunks their text, embeds the chunks,
and stores them persistently in ChromaDB.
For each question the app retrieves up to three relevant source chunks.

The app provides document Q&A, mock interviews and Python study.
Practice sessions use five questions, scores from 0 to 10, feedback and hints.

No production latency benchmark or employer deployment is documented here.
These notes describe public project features; they contain no private data.